

# Project

Victoria (Tori) Widjaja & Jeremiah Fa'atiligā

2024-06-19

## R Markdown

```
library(dplyr)
library(tidyr)
library(tibble)
library(readr)
library(tidyverse)
library(ggplot2)
library(GGally)
library(corrplot)
library(plotmo)
library(caret)
library(kernlab)
library(earth)
library(skimr)
library(psych)
library(reshape2)
library(gt)
library(ROSE)

rseed <- 123

dir_prefix <- getwd()
print(dir_prefix)

## [1] "C:/Users/ToriT/ADS503-Group6"

### Connection info for GitHub File
url <- paste(dir_prefix, 'healthcare-dataset-stroke-data.csv', sep = '/')
df_orig <- read_csv(url)
print(url)

## [1] "C:/Users/ToriT/ADS503-Group6/healthcare-dataset-stroke-data.csv"

describe(df_orig)
```

|                  | vars | n    | mean     | sd       | median   | trimmed  | mad      | min   |
|------------------|------|------|----------|----------|----------|----------|----------|-------|
| ## id            | 1    | 5110 | 36517.83 | 21161.72 | 36932.00 | 36542.26 | 27413.27 | 67.00 |
| ## gender*       | 2    | 5110 | 1.41     | 0.49     | 1.00     | 1.39     | 0.00     | 1.00  |
| ## age           | 3    | 5110 | 43.23    | 22.61    | 45.00    | 43.61    | 26.69    | 0.08  |
| ## hypertension  | 4    | 5110 | 0.10     | 0.30     | 0.00     | 0.00     | 0.00     | 0.00  |
| ## heart_disease | 5    | 5110 | 0.05     | 0.23     | 0.00     | 0.00     | 0.00     | 0.00  |
| ## ever_married* | 6    | 5110 | 1.66     | 0.48     | 2.00     | 1.70     | 0.00     | 1.00  |

```
## work_type*          7 5110      3.50      1.28      4.00      3.62      0.00      1.00
## Residence_type*     8 5110      1.51      0.50      2.00      1.51      0.00      1.00
## avg_glucose_level   9 5110    106.15    45.28    91.88    97.85    26.06    55.12
## bmi*               10 5110    172.19    88.96   158.00   163.08    74.13    1.00
## smoking_status*    11 5110      2.59      1.09      2.00      2.61      1.48    1.00
## stroke             12 5110      0.05      0.22      0.00      0.00      0.00    0.00
##
##           max      range  skew kurtosis      se
## id          72940.00 72873.00 -0.02   -1.21 296.03
## gender*         3.00      2.00  0.35   -1.86  0.01
## age           82.00    81.92 -0.14   -0.99  0.32
## hypertension     1.00      1.00  2.71    5.37  0.00
## heart_disease     1.00      1.00  3.94   13.57  0.00
## ever_married*     2.00      1.00 -0.66   -1.57  0.01
## work_type*        5.00      4.00 -0.91   -0.49  0.02
## Residence_type*    2.00      1.00 -0.03   -2.00  0.01
## avg_glucose_level 271.74    216.62  1.57    1.68  0.63
## bmi*            419.00    418.00  0.97    0.87  1.24
## smoking_status*    4.00      3.00  0.08   -1.35  0.02
## stroke           1.00      1.00  4.19   15.57  0.00
```

## Exploratory Data Analysis (EDA)

*###graphical and non-graphical representations of relationships between the response variable and predi*

```
df_eda <- df_orig
```

```
rownames(df_eda) <- df_eda$id
```

```
## Warning: Setting row names on a tibble is deprecated.
```

```
df_eda <- dplyr::select(df_eda, -id)
```

```
print(df_eda)
```

```
## # A tibble: 5,110 x 11
```

```
##   gender  age hypertension heart_disease ever_married work_type Residence_type
##   <chr>  <dbl>         <dbl>         <dbl> <chr>         <chr>      <chr>
## 1 Male    67           0           1 Yes      Private      Urban
## 2 Female  61           0           0 Yes      Self-emp~    Rural
## 3 Male    80           0           1 Yes      Private      Rural
## 4 Female  49           0           0 Yes      Private      Urban
## 5 Female  79           1           0 Yes      Self-emp~    Rural
## 6 Male    81           0           0 Yes      Private      Urban
## 7 Male    74           1           1 Yes      Private      Rural
## 8 Female  69           0           0 No        Private      Urban
## 9 Female  59           0           0 Yes      Private      Rural
## 10 Female 78           0           0 Yes      Private      Urban
```

```
## # i 5,100 more rows
```

```
## # i 4 more variables: avg_glucose_level <dbl>, bmi <chr>, smoking_status <chr>,
```

```
## #   stroke <dbl>
```

## Histograms

```

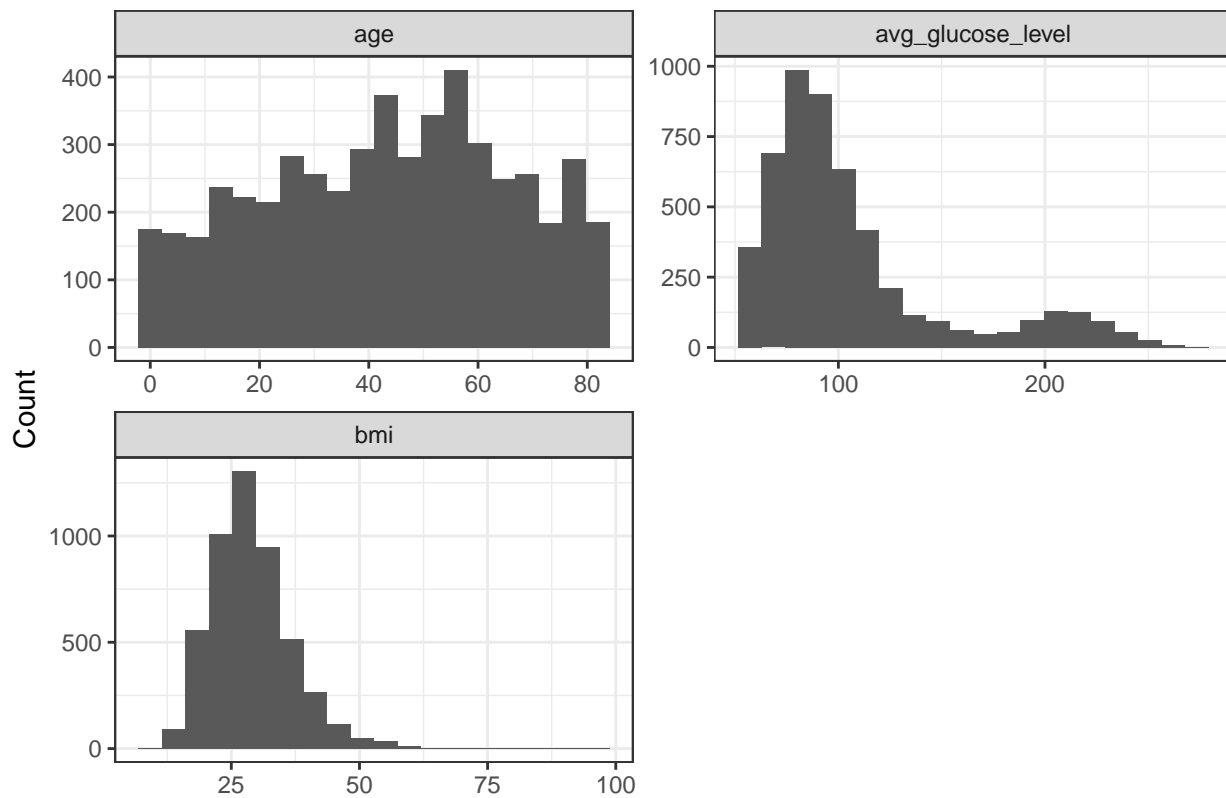
# Filter out N/A values for bmi and convert to numeric
df_eda <- df_eda %>%
  filter(!is.na(bmi)) %>%
  mutate(bmi = as.numeric(bmi))

# Pivot longer and convert value column to numeric if possible
df_long <- df_eda %>%
  pivot_longer(-c(stroke, ever_married, gender, hypertension, heart_disease, Residence_type, work_type),
  mutate(value = as.numeric(value))

# Plot histograms
ggplot(df_long, aes(x = value)) +
  geom_histogram(bins = 20) +
  facet_wrap(~Variable, scales = "free", ncol = 2) +
  theme_bw() +
  labs(title = 'Histograms of Variables Distributions', x = NULL, y = "Count")

```

Histograms of Variables Distributions



Corrplot

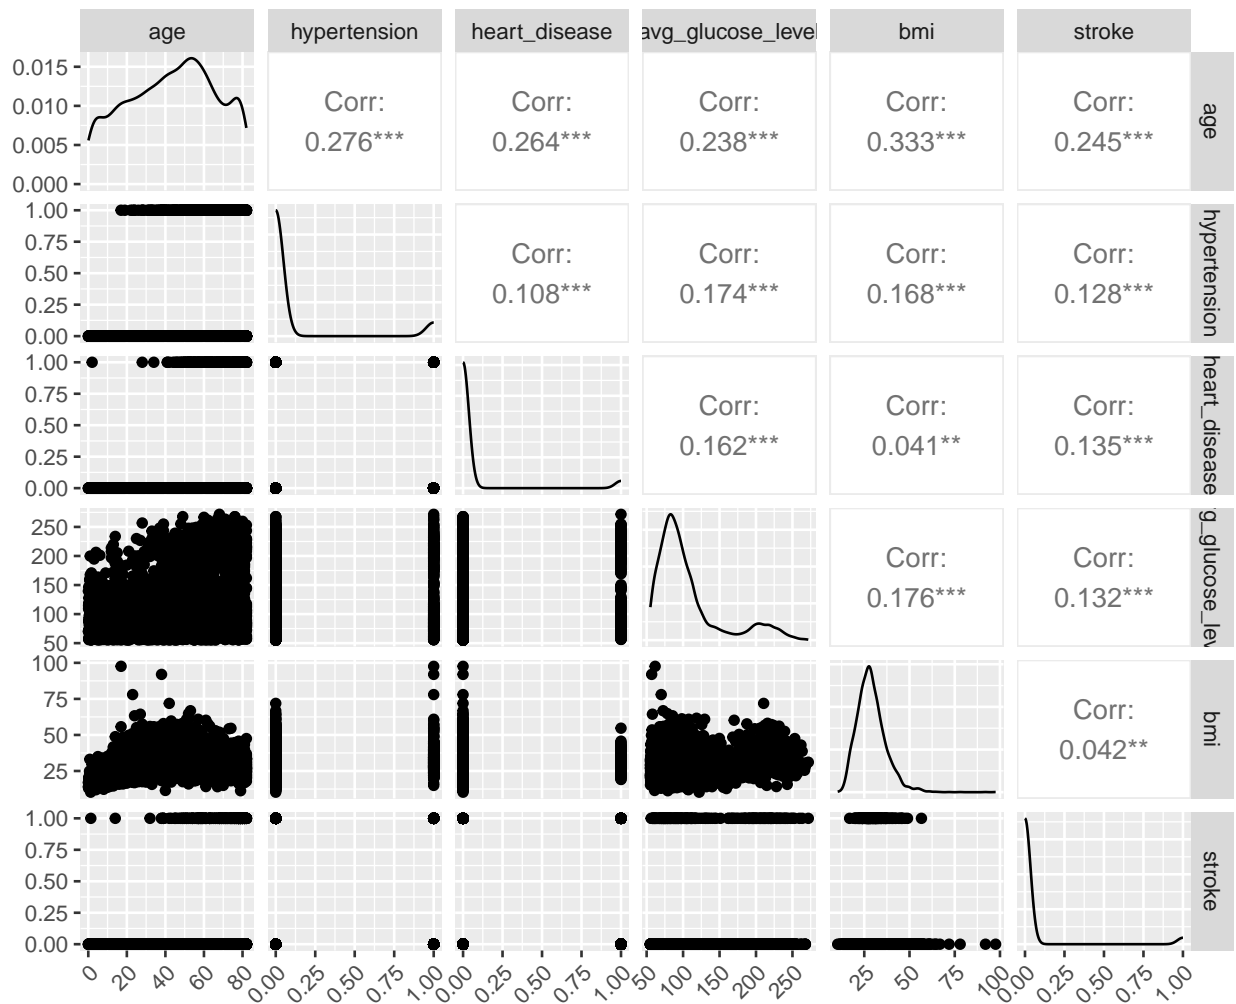
```

numeric_df <- df_eda[sapply(df_eda, is.numeric)]

numeric_df |>
  ggpairs(title = "Relationship Between Predictors", progress = FALSE)+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```

## Relationship Between Predictors



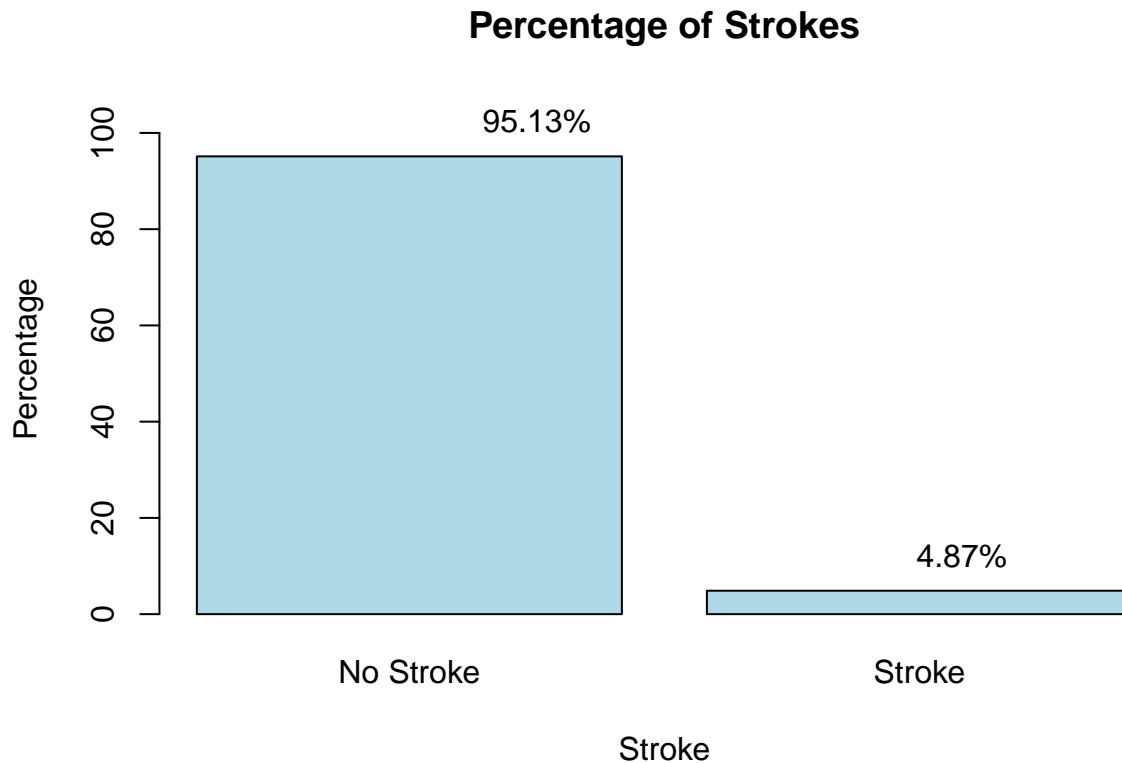
## Class Balance

```
stroke_percent <- prop.table(table(df_eda$stroke)) * 100

barplot(stroke_percent,
        main = "Percentage of Strokes",
        xlab = "Stroke",
        ylab = "Percentage",
        col = "lightblue",
        ylim = c(0, max(stroke_percent) + 11),
        names.arg = c("No Stroke", "Stroke"))

label_pos <- stroke_percent + 1

### Add labels
text(x = 1:length(stroke_percent),
     y = label_pos,
     labels = paste0(round(stroke_percent, 2), "%"),
     pos = 3)
```



#### Distributions of Values (Counts)

```
wrap_text <- function(x, width) {
  sapply(strwrap(x, width = width, simplify = FALSE), paste, collapse = "\n")
}

# Select non-numeric columns
non_numeric_columns <- names(df_eda)[sapply(df_eda, is.factor) | sapply(df_eda, is.character)]

# Function to create plots
create_plots <- function(col) {
  counts <- table(df_eda[[col]], df_eda$stroke)
  counts_df <- as.data.frame(counts)
  names(counts_df) <- c(col, "stroke", "count")

  ggplot(counts_df, aes(x = !!sym(col), y = count, fill = factor(stroke))) +
    geom_bar(stat = "identity", position = "dodge") +
    geom_text(aes(label = paste0(count),
      vjust = -0.5, size = 3, position = position_dodge(width = 0.9))) + # Add count / percent
    labs(x = col, y = "Count of Records",
      title = paste("Bar Graph of Stroke by", col, " by Count")) +
    theme(axis.text.x = element_text(angle = 45, hjust = 1),
      panel.spacing = unit(3, "lines"),
      legend.text = element_text(size = 7), # Adjust legend text size
      legend.title = element_text(size = 9)) # Adjust legend title size
}
```

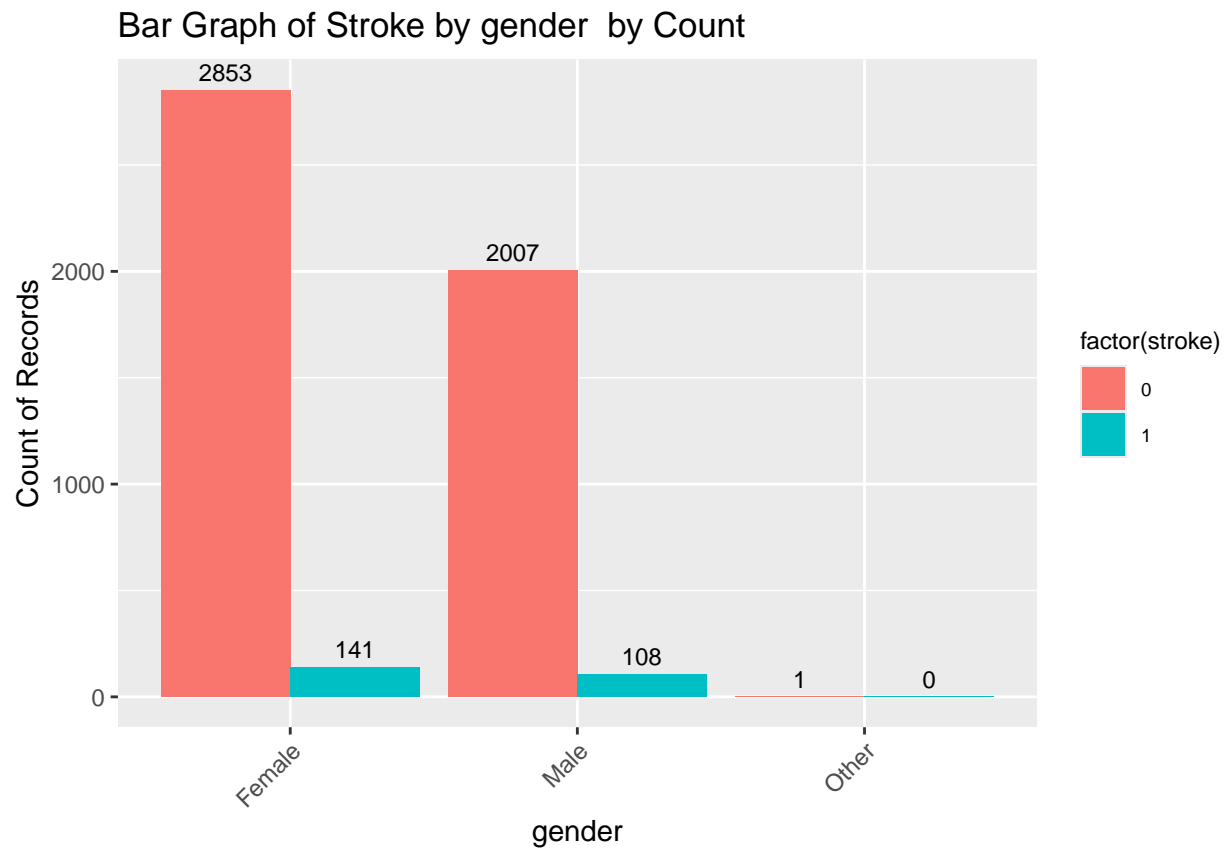
```

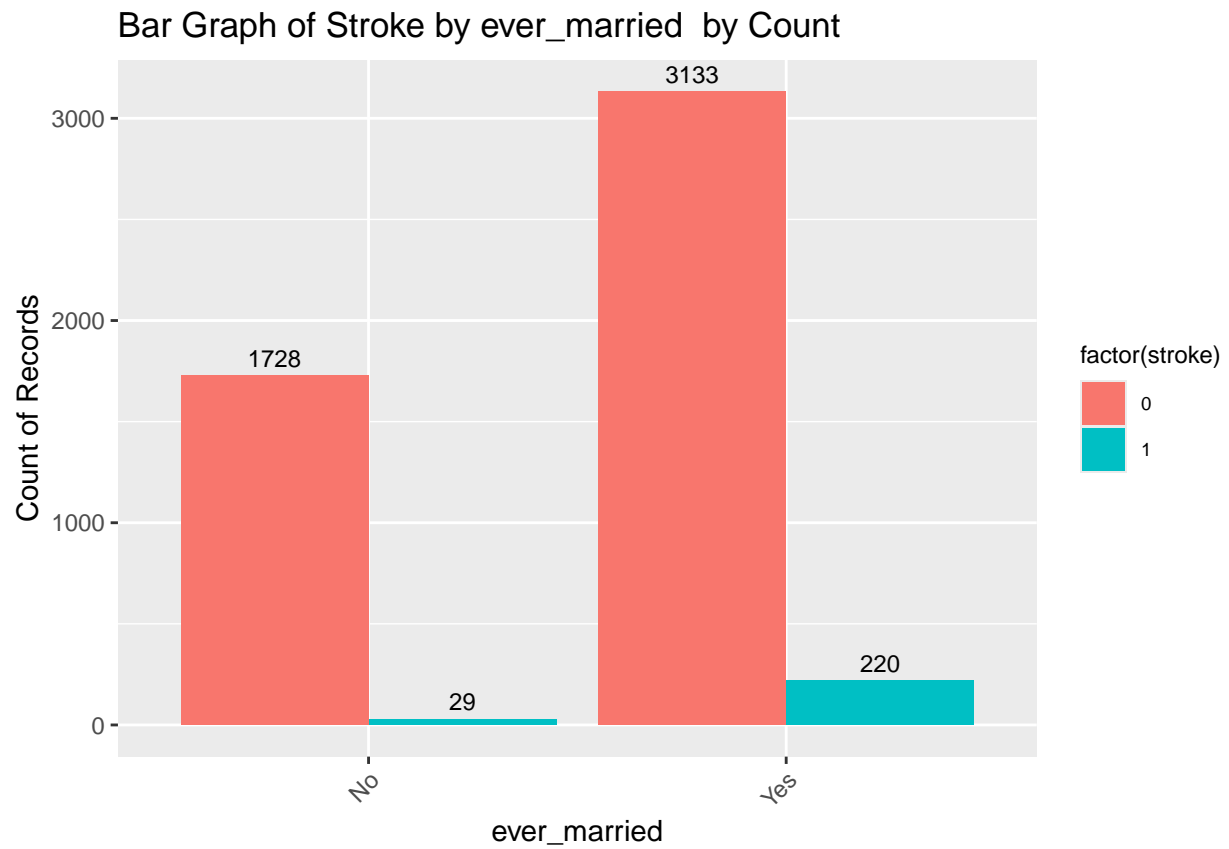
}

# Create plots for each non-numeric column
plots_list <- map(non_numeric_columns, ~ create_plots(.x))

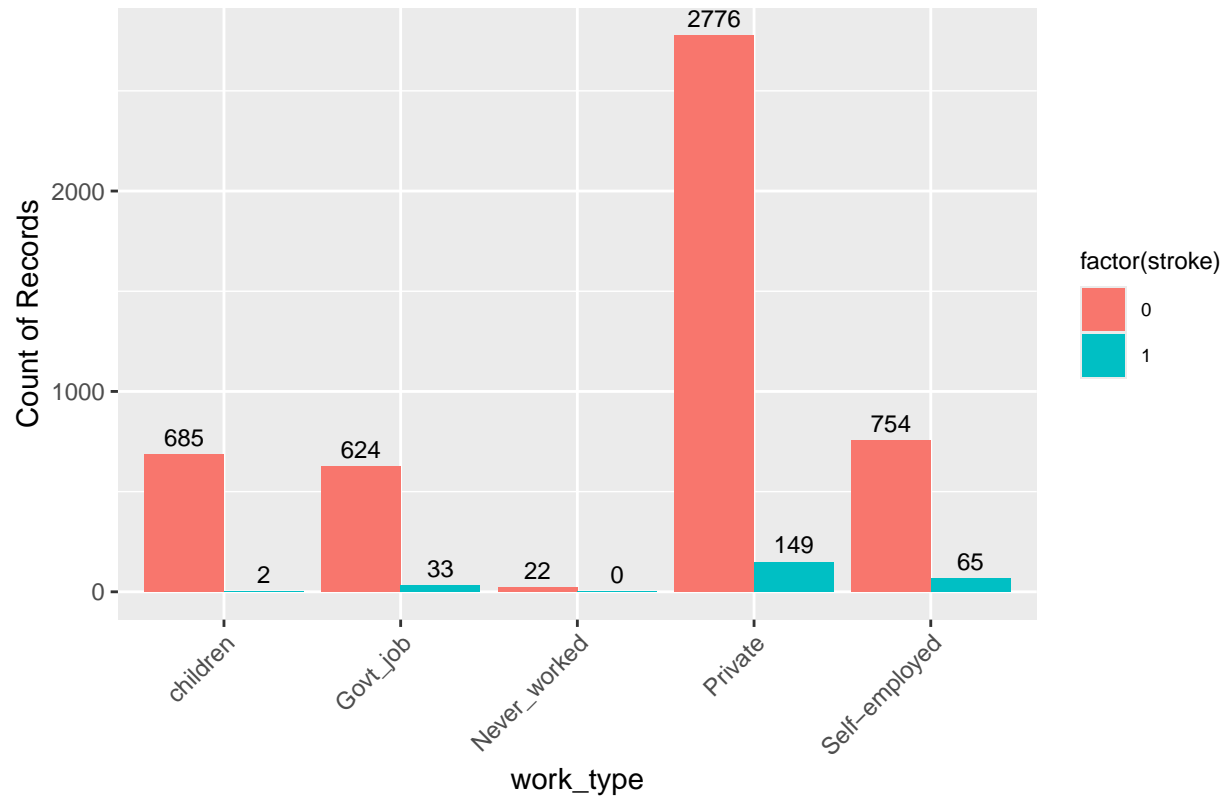
# Print plots
walk(plots_list, print)

```

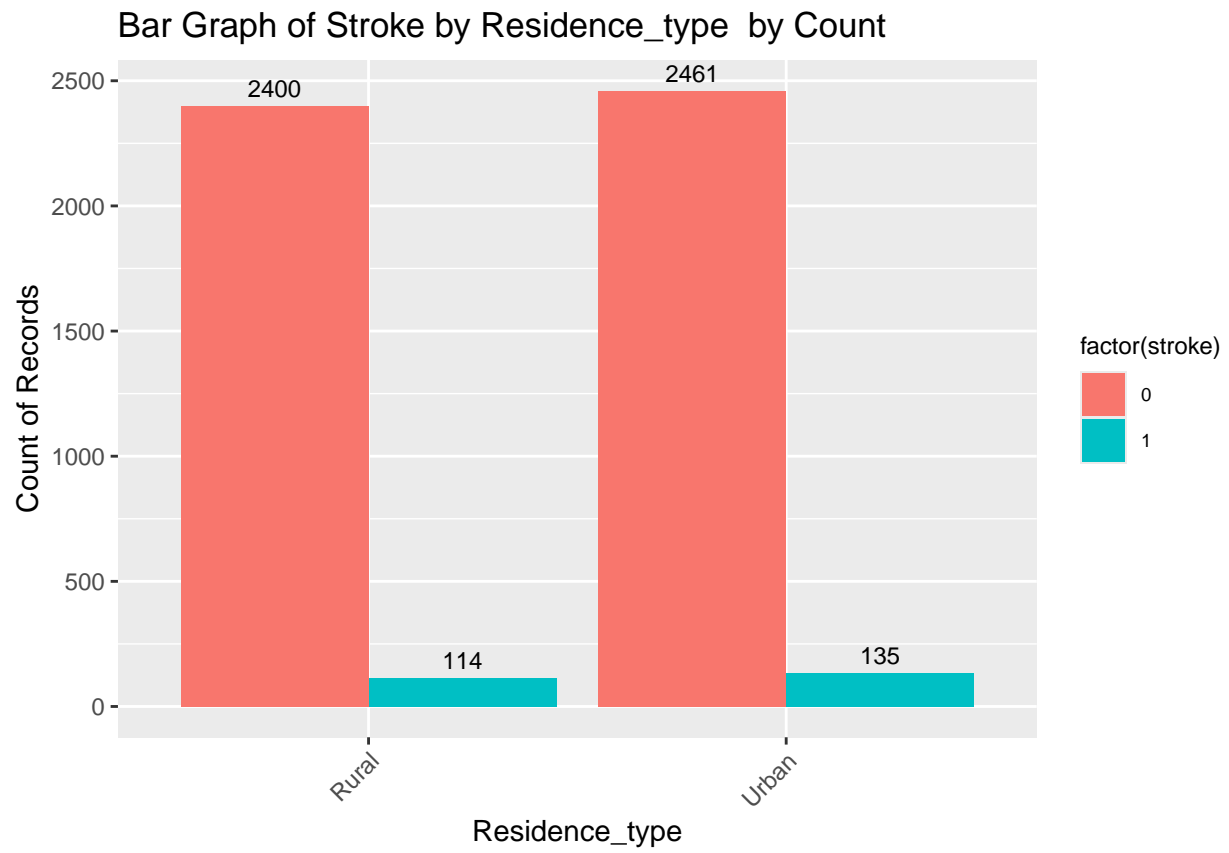


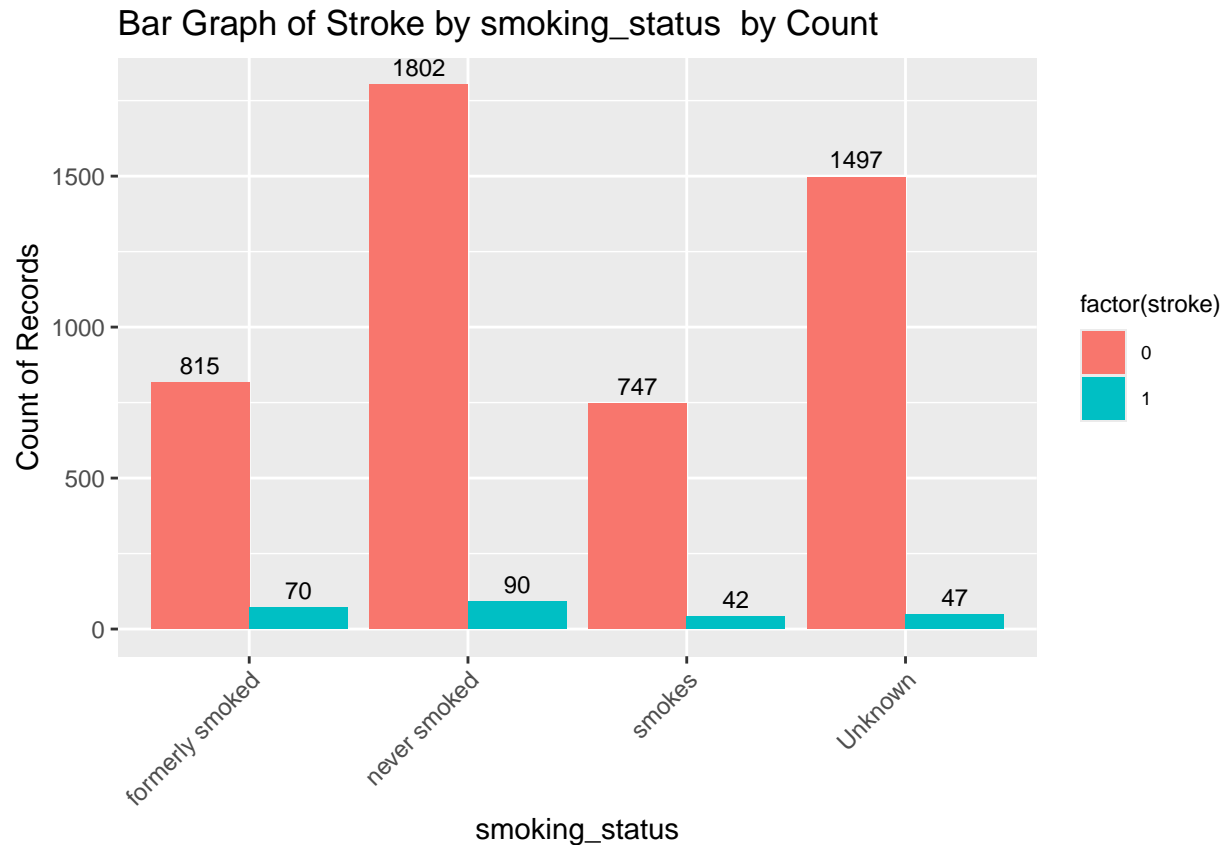


Bar Graph of Stroke by work\_type by Count









#### Distributions of Values (Percentages)

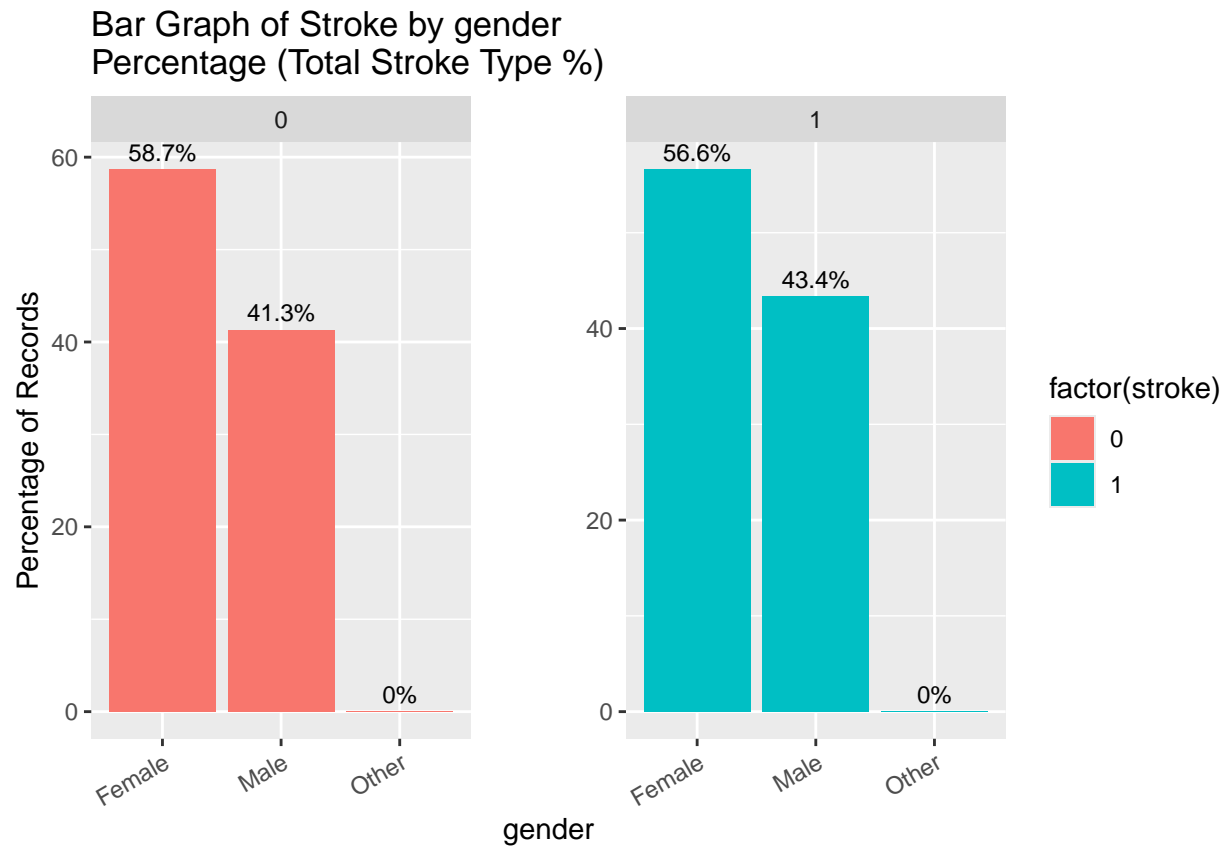
```
# Function to create plots
create_plots <- function(col) {
  counts <- table(df_eda[[col]], df_eda$stroke)
  counts_df <- as.data.frame(counts)
  names(counts_df) <- c(col, "stroke", "count")

  # Calculate percentages within each category
  counts_df <- counts_df %>%
    group_by(stroke) %>%
    mutate(percent = count / sum(count) * 100) %>%
    ungroup()

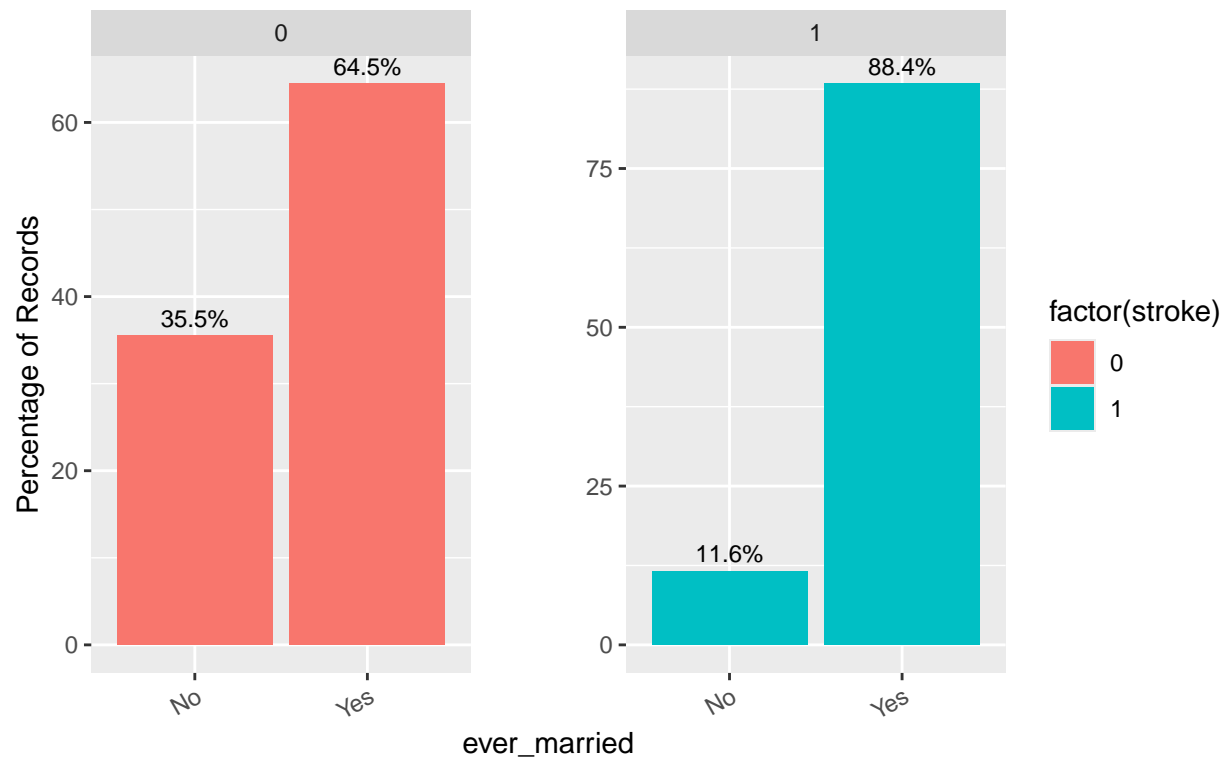
  ggplot(counts_df, aes(x = !!sym(col), y = percent, fill = factor(stroke))) +
    geom_bar(stat = "identity") +
    geom_text(aes(label = paste0(round(percent, 1), "%"),
      vjust = -0.5, size = 3) + # Add percent labels
    labs(x = col, y = "Percentage of Records",
      title = paste("Bar Graph of Stroke by", col, "\nPercentage (Total Stroke Type %)") +
    facet_wrap(~stroke, scales = "free") + # Facet by stroke
    theme(axis.text.x = element_text(angle = 30, hjust = 1),
      panel.spacing = unit(3, "lines"))
}
```

```
# Create plots for each non-numeric column
plots_list <- map(non_numeric_columns, ~ create_plots(.x))

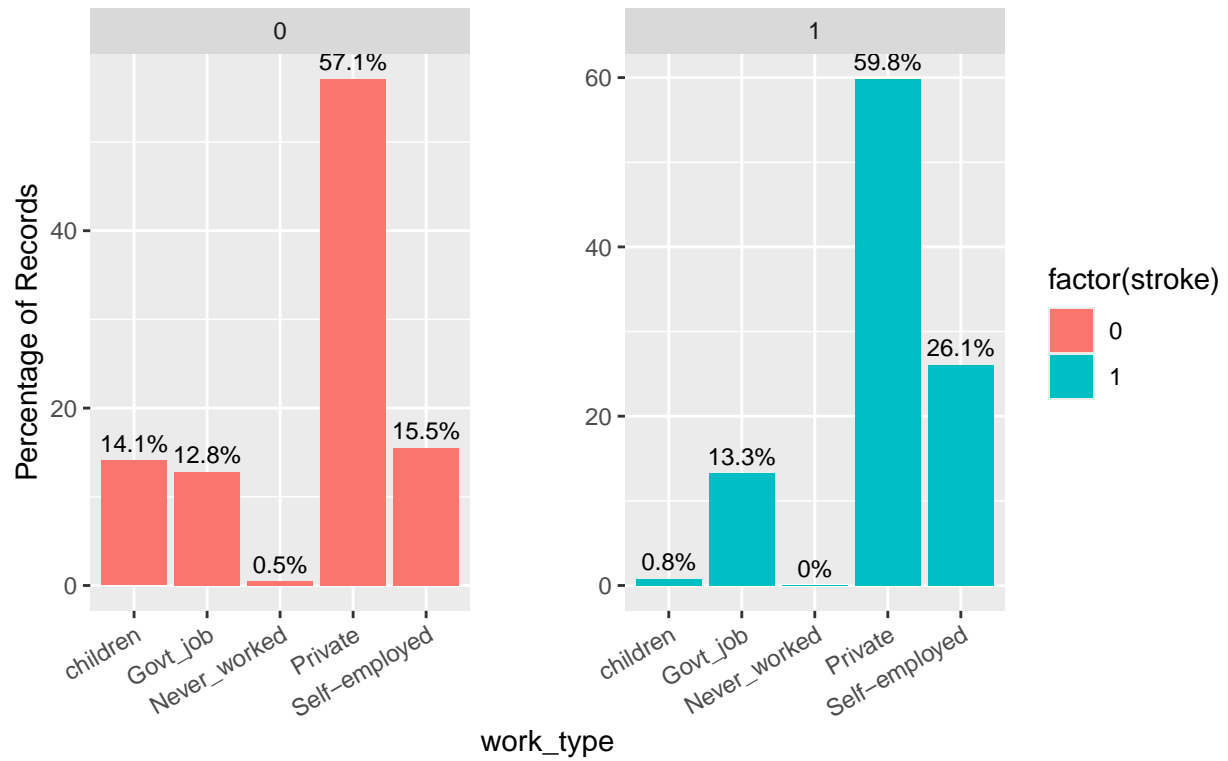
# Print plots
walk(plots_list, print)
```




Bar Graph of Stroke by ever\_married  
Percentage (Total Stroke Type %)




Bar Graph of Stroke by work\_type  
Percentage (Total Stroke Type %)





Bar Graph of Stroke by smoking\_status  
Percentage (Total Stroke Type %)

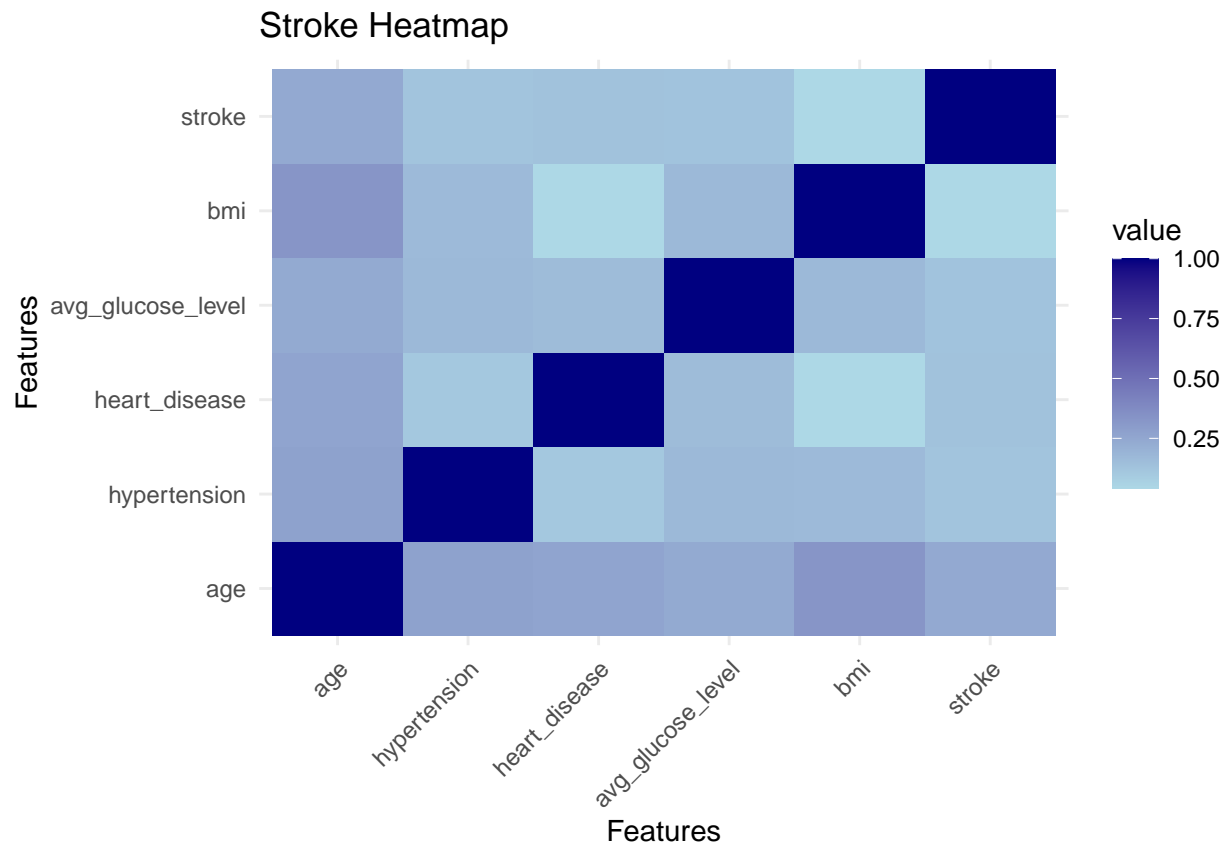


## Heatmap

```
##### HEAT MAP ##### For Feature Selection #####
# Numeric Columns for EDA
heatmap_data <- df_eda %>%
  select(age, hypertension, heart_disease, avg_glucose_level, bmi, stroke)

# Correlation Matrix
correlation_matrix <- cor(heatmap_data, use = "pairwise.complete.obs")

# Heatmap Plot
ggplot(data = melt(correlation_matrix), aes(x = Var1, y = Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradient(low = "Light Blue", high = "Navy Blue") +
  labs(x = "Features", y = "Features", title = "Stroke Heatmap") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



## Data Splitting - Training & Test

```
###training, validation, and test sets
###Split
trainIndex <- createDataPartition(df_eda$stroke, p = 0.8, ### saving 20% for test
                                  list = FALSE,
                                  times = 1)

# Subset data into training and testing sets
trainData <- df_eda[trainIndex, ]
testData <- df_eda[-trainIndex, ]
```

## Data Wrangling and Pre-Processing

### Missing Values

```
table(trainData$gender)

##
## Female   Male   Other
##   2395   1692     1

trainData$gender[trainData$gender == "Other"] <- {
  ux <- unique(trainData$gender)
  mode_gender <- ux[which.max(tabulate(match(trainData$gender, ux)))]
  mode_gender
```



```

}

testData$gender[testData$gender == "Other"] <- {
  ux <- unique(testData$gender)
  mode_gender <- ux[which.max(tabulate(match(testData$gender, ux)))]
  mode_gender
}
table(trainData$gender)

##
## Female    Male
##   2396    1692

# Check for missing data
colSums(is.na(trainData))

##           gender           age      hypertension      heart_disease
##             0             0             0             0
## ever_married      work_type  Residence_type avg_glucose_level
##             0             0             0             0
##           bmi      smoking_status           stroke
##          159             0             0

### bagImpute used because missing data is random
bag_missing <- preProcess(trainData, method = "bagImpute") ##### CHANGE? knn ??
trainData <- predict(bag_missing, newdata = trainData)

# bagImpute for Test Data
testData <- predict(bag_missing, newdata = testData)

colSums(is.na(trainData))

##           gender           age      hypertension      heart_disease
##             0             0             0             0
## ever_married      work_type  Residence_type avg_glucose_level
##             0             0             0             0
##           bmi      smoking_status           stroke
##            0             0             0

```

## Setting up Factors

```

### Factor Train
non_numeric_cols <- sapply(trainData, function(x) !is.numeric(x))
# Convert non-numeric columns to factors
trainData <- trainData %>%
  mutate_if(non_numeric_cols, as.factor)

### Factor Test
non_numeric_cols <- sapply(testData, function(x) !is.numeric(x))
# Convert non-numeric columns to factors Test
testData <- testData %>%
  mutate_if(non_numeric_cols, as.factor)

#str(trainData)
#str(testData)

```

## Dummy Variables

```
### dummy variables for Trfactors
dummy_model1 <- dummyVars(stroke ~ gender + ever_married + work_type + Residence_type + smoking_status,

trainData_dummy <- as.data.frame(predict(dummy_model1, newdata = trainData))
trainData_dummy <- as.data.frame(lapply(trainData_dummy, as.factor))

### dummy variables for test factors
dummy_model2 <- dummyVars(stroke ~ gender + ever_married + work_type + Residence_type + smoking_status,

testData_dummy <- as.data.frame(predict(dummy_model2, newdata = testData))
testData_dummy <- as.data.frame(lapply(testData_dummy, as.factor))

#head(trainData_dummy)
#head(testData_dummy)

### Dropping columns that can be inferred from the others to avoid multicollinearity
trainData_dummy <-select(trainData_dummy, -gender.Female, -ever_married.No, -work_type.children, -Residence_type.Children)

#head(trainData_dummy)
trainData_selected <- select(trainData, age, hypertension, heart_disease, avg_glucose_level, bmi, stroke)
#str(trainData_selected)

### Test Dropping columns that can be inferred from the others to avoid multicollinearity
testData_dummy <-select(testData_dummy, -gender.Female, -ever_married.No, -work_type.children, -Residence_type.Children)
testData_selected <- select(testData, age, hypertension, heart_disease, avg_glucose_level, bmi, stroke)
#str(testData_dummy)

# Train
trainData_ready <- cbind(trainData_dummy, trainData_selected)
trainData_ready <- trainData_ready %>%
  mutate(across(all_of(c("hypertension", "heart_disease", "stroke")), as.factor))
names(trainData_ready) <- gsub("\\\\.", "_", names(trainData_ready))
names(trainData_ready) <- gsub("_Yes", "", names(trainData_ready))
trainData_ready$stroke <- as.factor(make.names(as.character(trainData_ready$stroke)))
head(trainData_ready)

##   gender_Male ever_married work_type_Govt_job work_type_Never_worked
## 1           1           1                   0                      0
## 2           0           1                   0                      0
## 3           0           1                   0                      0
## 4           1           1                   0                      0
## 5           1           1                   0                      0
## 6           0           0                   0                      0
##   work_type_Private work_type_Self-employed Residence_type_Urban
## 1                 1                 0                 1
## 2                 0                 1                 0
## 3                 1                 0                 1
## 4                 1                 0                 1
## 5                 1                 0                 0
## 6                 1                 0                 1
##   smoking_status_formerly_smoked smoking_status_never_smoked
## 1                             1                             0
## 2                             0                             1
```

```
## 3          0          0
## 4          1          0
## 5          0          1
## 6          0          1
##   smoking_status_smokes age hypertension heart_disease avg_glucose_level
## 1          0 67          0          1          228.69
## 2          0 61          0          0          202.21
## 3          1 49          0          0          171.23
## 4          0 81          0          0          186.21
## 5          0 74          1          1          70.09
## 6          0 69          0          0          94.39
##           bmi stroke
## 1 36.60000    X1
## 2 33.74876    X1
## 3 34.40000    X1
## 4 29.00000    X1
## 5 27.40000    X1
## 6 22.80000    X1
```

```
# Test
testData_ready <- cbind(testData_dummy, testData_selected)
testData_ready <- testData_ready %>%
  mutate(across(all_of(c("hypertension", "heart_disease", "stroke")), as.factor))
names(testData_ready) <- gsub("\\.", "_", names(testData_ready))
names(testData_ready) <- gsub("_Yes", "", names(testData_ready))
testData_ready$stroke <- as.factor(make.names(as.character(testData_ready$stroke)))
str(testData_ready)
```


```
## 'data.frame':   1022 obs. of  16 variables:
## $ gender_Male      : Factor w/ 2 levels "0","1": 2 1 1 1 1 2 2 2 1 2 ...
## $ ever_married     : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 1 2 2 ...
## $ work_type_Govt_job : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 2 1 1 ...
## $ work_type_Never_worked : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ work_type_Private : Factor w/ 2 levels "0","1": 2 1 2 1 1 1 1 1 2 2 ...
## $ work_type_Self-employed : Factor w/ 2 levels "0","1": 1 2 1 2 2 2 2 1 1 1 ...
## $ Residence_type_Urban : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 2 2 2 1 ...
## $ smoking_status_formerly_smoked: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 2 1 2 1 ...
## $ smoking_status_never_smoked : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 1 2 1 2 ...
## $ smoking_status_smokes : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ age              : num  80 79 81 50 52 80 80 48 63 76 ...
## $ hypertension     : Factor w/ 2 levels "0","1": 1 2 2 2 2 1 1 1 1 2 ...
## $ heart_disease     : Factor w/ 2 levels "0","1": 2 1 1 1 1 1 2 1 1 1 ...
## $ avg_glucose_level : num  105.9 174.1 80.4 167.4 233.3 ...
## $ bmi              : num  32.5 24 29.7 30.9 48.9 ...
## $ stroke            : Factor w/ 2 levels "X0","X1": 2 2 2 2 2 2 2 2 2 2 ...
```

### Class Imbalance (ROSE Method)

```
# Calculate counts and convert to data frame
stroke_counts <- as.data.frame(table(trainData_ready$stroke))
colnames(stroke_counts) <- c("Stroke", "Count")
stroke_counts$Stroke <- factor(stroke_counts$Stroke, levels = c('X0', 'X1'), labels = c("No Stroke", "S

### Bar Chart
ggplot(stroke_counts, aes(x = Stroke, y = Count)) +
```

```
geom_bar(stat = "identity", fill = "skyblue") +
labs(title = "Count of Stroke Values",
      x = "Stroke",
      y = "Count") +
geom_text(aes(label = Count),
          vjust = -0.5, # Position above the bars
          size = 3, # Size of the text
          color = "black") # Color of the text
```



```
set.seed(rseed)

# Apply ROSE to balance the dataset
rose_train <- ROSE(stroke ~ ., data = trainData_ready)$data

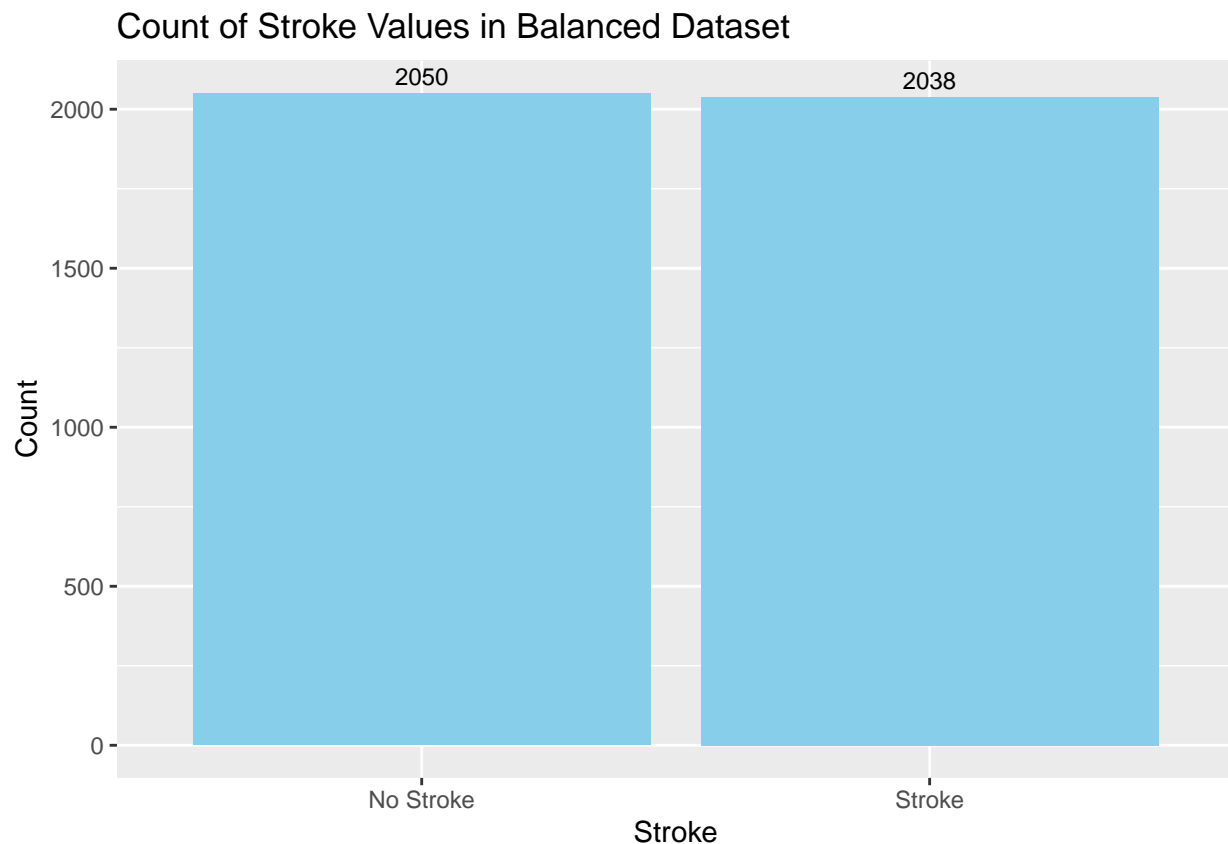
# Separate predictors (features) and target variable (stroke)
train_X <- rose_train[, !(names(rose_train) %in% "stroke")]
train_y <- rose_train$stroke

train_y_df <- data.frame(stroke = train_y)

stroke_counts <- as.data.frame(table(train_y_df$stroke))
colnames(stroke_counts) <- c("Stroke", "Count")
stroke_counts$Stroke <- factor(stroke_counts$Stroke,
                              levels = c('X0', 'X1'), labels = c("No Stroke", "Stroke"))

### Bar chart
ggplot(stroke_counts, aes(x = Stroke, y = Count)) +
```

```
geom_bar(stat = "identity", fill = "skyblue") +
labs(title = "Count of Stroke Values in Balanced Dataset",
     x = "Stroke",
     y = "Count") +
geom_text(aes(label = Count),
          vjust = -0.5,
          size = 3,
          color = "black")
```



```
trainData_ready <- data.frame(rose_train)
```

### Baseline Model to Beat

```
set.seed(123)
### Parameters for Tuning
ctrl <- trainControl(method = "repeatedcv",
                     number = 5,
                     repeats = 3,
                     verboseIter = FALSE,
                     classProbs = TRUE,
                     summaryFunction = twoClassSummary)

### Base Model Logistic Regression Model

# Fit logistic regression model
lrm_base_model <- train(stroke ~ .,
```

```

data = trainData_ready,
method = "glm",
family = "binomial",
trControl = ctrl)

## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.

# View model summary
summary(lrm_base_model)

##
## Call:
## NULL
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.660e+00  2.903e-01 -12.606  < 2e-16 ***
## gender_Male1    -3.688e-01  8.355e-02  -4.414  1.01e-05 ***
## ever_married1    3.268e-01  1.253e-01   2.607  0.009126 **
## work_type_Govt_job1 -1.269e+00  3.323e-01  -3.819  0.000134 ***
## work_type_Never_worked1 -1.137e+01  1.879e+02  -0.060  0.951759
## work_type_Private1 -9.997e-01  3.199e-01  -3.125  0.001780 **
## work_type_Self-employed1 -1.320e+00  3.367e-01  -3.921  8.82e-05 ***
## Residence_type_Urban1  2.210e-01  7.961e-02   2.776  0.005505 **
## smoking_status_formerly_smoked1  4.485e-01  1.240e-01   3.618  0.000296 ***
## smoking_status_never_smoked1  -2.937e-02  1.139e-01  -0.258  0.796570
## smoking_status_smokes1  4.335e-01  1.308e-01   3.314  0.000919 ***
## age             6.697e-02  2.854e-03  23.461  < 2e-16 ***
## hypertension1    4.910e-01  1.062e-01   4.623  3.79e-06 ***
## heart_disease1    6.163e-01  1.317e-01   4.680  2.87e-06 ***
## avg_glucose_level  5.250e-03  7.211e-04   7.280  3.33e-13 ***
## bmi             -8.352e-03  5.905e-03  -1.414  0.157245
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 5667.1  on 4087  degrees of freedom
## Residual deviance: 3946.2  on 4072  degrees of freedom
## AIC: 3978.2
##
## Number of Fisher Scoring iterations: 12

```

## Data Prep

```

# Exclude stroke and Identify and remove variables with zero variance
Zero_Var_vars <- nearZeroVar(trainData_ready[, -which(names(trainData_ready) == "stroke"), drop = FALSE]
trainData_ready <- trainData_ready[, -Zero_Var_vars, drop = FALSE]
print(names(Zero_Var_vars))

## NULL

```

## Models

### Model #1 - Linear Discriminant Analysis (LDR)

```
rseed <- 123
set.seed(rseed)
# Linear Discriminant Analysis
ldaFit_stroke <- train(stroke ~ .,
                       data = trainData_ready,
                       method = 'lda',
                       preProc = c("center", "scale"),
                       metric = 'ROC',
                       trControl = ctrl)
ldaFit_stroke

## Linear Discriminant Analysis
##
## 4088 samples
## 14 predictor
## 2 classes: 'X0', 'X1'
##
## Pre-processing: centered (14), scaled (14)
## Resampling: Cross-Validated (5 fold, repeated 3 times)
## Summary of sample sizes: 3270, 3270, 3270, 3271, 3271, 3271, ...
## Resampling results:
##
## ROC      Sens      Spec
## 0.8426977 0.7256911 0.810274

ldaCM_stroke <- confusionMatrix(ldaFit_stroke, norm = "none")
ldaCM_stroke

## Cross-Validated (5 fold, repeated 3 times) Confusion Matrix
##
## (entries are un-normalized aggregated counts)
##
##           Reference
## Prediction  X0  X1
##           X0 4463 1160
##           X1 1687 4954
##
## Accuracy (average) : 0.7679
```

### Model #2 - Penalized Logistic Regression (PLR)

```
set.seed(rseed)

glmGrid <- expand.grid(alpha = c( 0, .5, 1),
                      lambda = seq(.001, 2, length = 20))


# Penalized Logistic Regression
plrFit_stroke <- train(stroke ~ .,
                      data = trainData_ready,
                      method = 'glmnet',
                      tuneGrid = glmGrid,
```

```

preProc = c('center', 'scale'),
metric = 'ROC',
trControl = ctrl)

plot(plrFit_stroke)

```



```

# Best tuning parameters
optimal_plr_tune <- plrFit_stroke$bestTune
print(paste('Best Alpha and Lambda tuning parameters for Penalized Logistic Regression:', paste(optimal_plr_tune)))

## [1] "Best Alpha and Lambda tuning parameters for Penalized Logistic Regression: 0.5,0.001"

```

### Model #3 - Nearest Shrunken Centroids

```

set.seed(rseed)
nscGrid <- data.frame(threshold = seq(2, 7, length = 20))

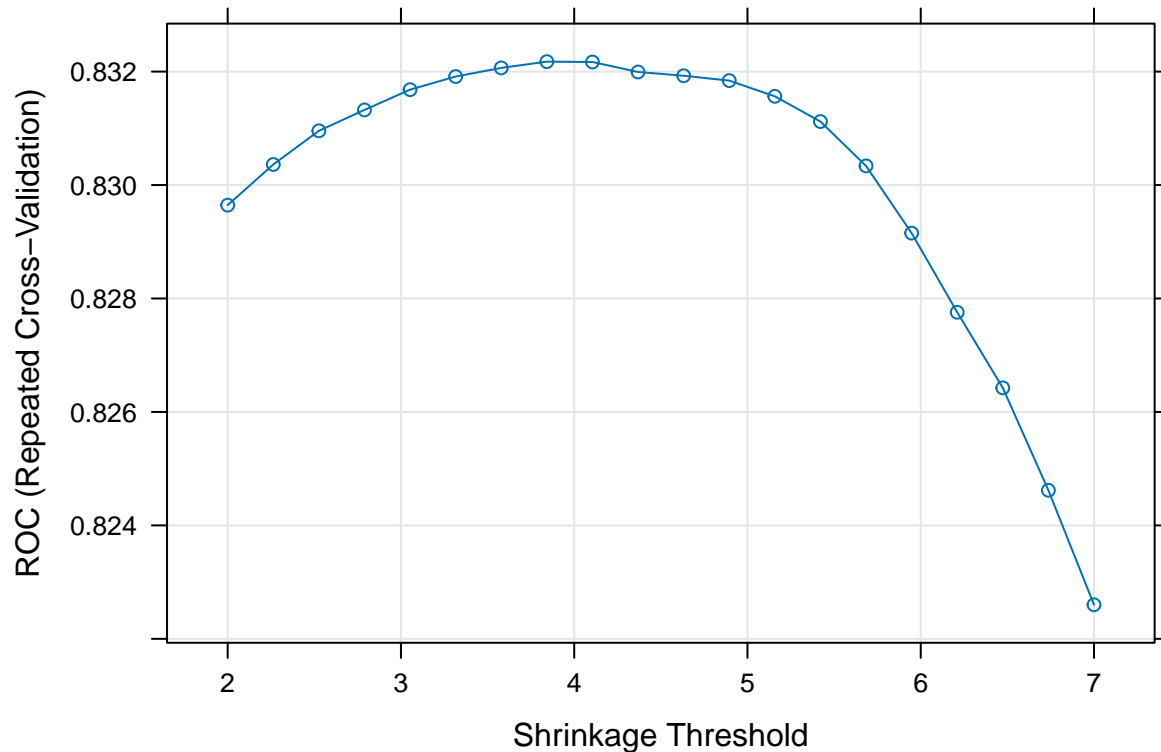
# Nearest Shrunken Centroids
nscFit_stroke <- train(stroke ~ .,
  data = trainData_ready,
  method = 'pam',
  preProc = c('center', 'scale'),
  tuneGrid = nscGrid,
  metric = 'ROC',
  trControl = ctrl)

```



```
## 111111111111111111
```

```
plot(nscFit_stroke)
```



```
# Best tuning parameters
```

```
optimal_nsc_tune <- nscFit_stroke$bestTune
```

```
print(paste('Best threshold tuning parameter for Nearest Shrunk Centroids:', paste(optimal_nsc_tune, 0)))
```

```
## [1] "Best threshold tuning parameter for Nearest Shrunk Centroids: 3.84210526315789"
```

#### Model #4 - Neural Network


```
set.seed(rseed)
```

```
grid <- expand.grid(size = c(1, 5, 8),      # Number of hidden units
                   decay = c( 0.01, 0.05, .1)) # Weight decay
```

```
# Neural Networks
```

```
nnFit_stroke <- train(stroke ~ .,
                     data = trainData_ready,
                     method = 'nnet',
                     maxit = 500,
                     preProcess = c('center', 'scale'),
                     metric = 'ROC',
                     trControl = ctrl,
                     tuneGrid = grid,
                     trace = FALSE)
```

```
plot(nnFit_stroke)
```



```
# Best tuning parameters
optimal_nn_tune <- nnFit_stroke$bestTune
print(paste('Best Size and Decay tuning parameters for Neural Network:', paste(optimal_nn_tune, collapse=', ')))

## [1] "Best Size and Decay tuning parameters for Neural Network: 8,0.1"
```

#### Model #5 - Random Forest


```
# Set seed for reproducibility
set.seed(rseed)

# Random Forest
rf_grid <- expand.grid(
  mtry = c(2, 4, 6)
)

rfFit_stroke <- train(
  stroke ~ .,
  data = trainData_ready,
  method = "rf",
  trControl = ctrl,
  tuneGrid = rf_grid,
  ntree = 100,
  trace = FALSE
)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
```

```
## in the result set. ROC will be used instead.
# Print the best parameters and performance metrics
plot(rfFit_stroke)
```



```
# Best tuning parameters
optimal_rf_tune <- rfFit_stroke$bestTune
print(paste('Best Size and Decay tuning parameters for Random Forest:', paste(optimal_rf_tune, collapse = ' ')))

## [1] "Best Size and Decay tuning parameters for Random Forest: 4"
```

## Results - Summary Table

```
# 1.Make predictions for Linear Discriminant Analysis
lda_pred <- predict(ldaFit_stroke, newdata = testData_ready)
lda_cm <- confusionMatrix(lda_pred, testData_ready$stroke)

# 2.Make predictions for Penalized Logistic Regression
plr_pred <- predict(plrFit_stroke, newdata = testData_ready)
plr_cm <- confusionMatrix(plr_pred, testData_ready$stroke)

# 3.Make predictions for Nearest Shrunken Centroids
nsc_pred <- predict(nscFit_stroke, newdata = testData_ready)
nsc_cm <- confusionMatrix(nsc_pred, testData_ready$stroke)

# 4.Make predictions for Neural Network
nn_pred <- predict(nnFit_stroke, newdata = testData_ready)
```

```

nn_cm <- confusionMatrix(nn_pred, testData_ready$stroke)

# 5. Make predictions for Random Forest
rf_pred <- predict(rfFit_stroke, newdata = testData_ready)
rf_cm <- confusionMatrix(rf_pred, testData_ready$stroke)

extract_accuracy <- function(model, cm) {
  tibble(
    model = model,
    Accuracy = cm$overall['Accuracy'],
    `CI Lower` = cm$overall['AccuracyLower'],
    `CI Upper` = cm$overall['AccuracyUpper']
  )
}

# Combine accuracies into a table
accuracies <- bind_rows(
  extract_accuracy('Linear Discriminant Analysis', lda_cm),
  extract_accuracy('Penalized Logistic Regression', plr_cm),
  extract_accuracy('Nearest Shrunken Centroids', nsc_cm),
  extract_accuracy('Neural Network', nn_cm),
  extract_accuracy('Random Forest', rf_cm)
)

# Display the table using gt package
accuracies %>%
  arrange(-Accuracy) %>%
  gt() %>%
  fmt_number(columns = c(Accuracy, `CI Lower`, `CI Upper`), decimals = 3)

```

| model                         | Accuracy | CI Lower | CI Upper |
|-------------------------------|----------|----------|----------|
| Random Forest                 | 0.779    | 0.752    | 0.804    |
| Neural Network                | 0.763    | 0.736    | 0.789    |
| Penalized Logistic Regression | 0.755    | 0.728    | 0.781    |
| Linear Discriminant Analysis  | 0.747    | 0.719    | 0.773    |
| Nearest Shrunken Centroids    | 0.710    | 0.681    | 0.738    |

## Tuned Models

```

set.seed(rseed)
# Penalized Logistic Regression with optimal parameters
plrFit_final <- train(stroke ~ .,
  data = trainData_ready,
  method = 'glmnet',
  preProc = c('center', 'scale'),
  metric = 'ROC',
  trControl = ctrl,
  tuneGrid = expand.grid(
    alpha = optimal_plr_tune$alpha,
    lambda = optimal_plr_tune$lambda
  ))
set.seed(rseed)

```

```

# Retrain NSC with optimal parameters
nscFit_final <- train(stroke ~ .,
                      data = trainData_ready,
                      method = 'pam',
                      preProc = c('center', 'scale'),
                      metric = 'ROC',
                      trControl = ctrl,
                      tuneGrid = expand.grid(
                        threshold = optimal_nsc_tune$threshold
                      ))

```

```
## 111111111111111111
```

```

set.seed(rseed)
# Retrain NN with optimal parameters
nnFit_final <- train(stroke ~ .,
                    data = trainData_ready,
                    method = 'nnet',
                    maxit = 500,
                    preProcess = c('center', 'scale'),
                    metric = 'ROC',
                    trControl = ctrl,
                    tuneGrid = expand.grid(
                      size = optimal_nn_tune$size,
                      decay = optimal_nn_tune$decay
                    ),
                    trace = FALSE)

set.seed(rseed)
# Retrain RF with optimal parameters
rfFit_final <- train(stroke ~ .,
                    data = trainData_ready,
                    method = "rf",
                    ntree = 100,
                    trControl = ctrl,
                    tuneGrid = expand.grid(
                      mtry = optimal_rf_tune$mtry
                    ))

```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

## Summary for Tuned Models

```

set.seed(rseed)
# Make predictions for Linear Discriminant Analysis
lda_pred <- predict(ldaFit_stroke, newdata = testData_ready)
lda_cm2 <- confusionMatrix(lda_pred, testData_ready$stroke)

# Make predictions for Penalized Logistic Regression
plr_pred2 <- predict(plrFit_final, newdata = testData_ready)
plr_cm2 <- confusionMatrix(plr_pred2, testData_ready$stroke)

# Make predictions for Nearest Shrunken Centroids
nsc_pred2 <- predict(nscFit_final, newdata = testData_ready)

```

```

nsc_cm2 <- confusionMatrix(nsc_pred2, testData_ready$stroke)

# Make predictions for Neural Network
nn_pred2 <- predict(nnFit_final, newdata = testData_ready)
nn_cm2 <- confusionMatrix(nn_pred2, testData_ready$stroke)

# Make predictions for Random Forest
rf_pred2 <- predict(rfFit_final, newdata = testData_ready)
rf_cm2 <- confusionMatrix(rf_pred2, testData_ready$stroke)

extract_accuracy <- function(model, cm) {
  tibble(
    model = model,
    Accuracy = cm$overall['Accuracy'],
    `CI Lower` = cm$overall['AccuracyLower'],
    `CI Upper` = cm$overall['AccuracyUpper']
  )
}

# Combine accuracies into a table
accuracies <- bind_rows(
  extract_accuracy('Linear Discriminant Analysis', lda_cm2),
  extract_accuracy('Penalized Logistic Regression', plr_cm2),
  extract_accuracy('Nearest Shrunken Centroids', nsc_cm2),
  extract_accuracy('Neural Network', nn_cm2),
  extract_accuracy('Random Forest', rf_cm2),
)

# Display the table using gt package
accuracies %>%
  arrange(-Accuracy) %>%
  gt() %>%
  fmt_number(columns = c(Accuracy, `CI Lower`, `CI Upper`), decimals = 3)

```

| model                         | Accuracy | CI Lower | CI Upper |
|-------------------------------|----------|----------|----------|
| Random Forest                 | 0.771    | 0.744    | 0.796    |
| Neural Network                | 0.763    | 0.736    | 0.789    |
| Penalized Logistic Regression | 0.755    | 0.728    | 0.781    |
| Linear Discriminant Analysis  | 0.747    | 0.719    | 0.773    |
| Nearest Shrunken Centroids    | 0.710    | 0.681    | 0.738    |